

FIRST: Stale Rollouts Can Pick the Wrong Training Data

A Measurement-First Audit of Off-Policy Data Selection in RLVR

Anonymous*

Abstract

Data selection for reinforcement learning with verifiable rewards (RLVR) studies the problem of choosing which prompts to train on. Because fresh rollouts for every candidate prompt are expensive, existing methods rank prompts with *stale* rollouts generated by an earlier policy, corrected by an importance weight over only one part of the trajectory. The uncorrected part introduces a bias that current practice assumes is negligible, but whether the resulting rankings can actually be trusted has not been examined. This raises two questions: (I) When can a stale ranking be trusted? (II) What should replace it when it cannot? For (I), we prove that any one-sided correction can flip the sign of a prompt’s gradient while every monitored training metric stays normal, and we show empirically that whether this failure occurs is decided by one measurable quantity: the reliability of the ranking signal in the prompt pool, which we estimate with a split-half floor. For (II), we propose FIRST (Floor-Informed Reliability of SelecTion), a procedure that reads the artifacts a training run already produces and prescribes, with no new rollouts, one of three policies: uniform sampling, a zero-cost pass-rate ranking, or stale scores with the disagreeing prompts sent to a small audit. On every pool we evaluated, the pre-registered verdict was that no recoverable ranking signal exists, and the zero-cost heuristic matched every gradient-based method. Together, these results suggest that the reliability of the ranking signal should be measured before any selection method is applied or evaluated.

1 Introduction

RLVR pipelines choose their training data by scoring each candidate prompt with “how much would training on it help the current policy π ?” and keeping the top k . Exact scores need fresh rollouts from π for every candidate, at every selection round. That is expensive. So practice reuses the rollouts and rewards produced by an earlier behavior policy β , and applies an importance correction. The correction has two axes: (a) how the current policy *reaches* a token (prefix occupancy), and (b) how it *finishes* afterwards (continuation outcome). Existing estimators correct exactly one axis. This is not an oversight: correcting both multiplies hundreds of token ratios, and the variance explodes. The implicit assumption is that the uncorrected half adds only a small bias—small enough to leave the *ranking* intact.

Published evidence already challenges this assumption. The method that popularized single-token influence scoring reports, in its own appendix, two relevant numbers: on- and off-policy per-prompt gradients agree pointwise (cosine above 0.6 on 40 of 50 prompts), yet only 28.8% of their top-10% neighborhood structure survives [CROPI]. We treat this as a symptom, not as

*Draft v0.5 (2026-08-18). Reframed structure following the pre-registered “structural absence” branch; numbers are from completed runs (v1 gate + v2 corrected, 2 seeds); pending seed harvests will update Sections 5 and 8 only.

proof of our mechanism. The number is not decomposed over the two correction axes, and other explanations—validation noise, difficulty clustering, score calibration—can damage rank structure too. This paper supplies what is missing: the mechanism class that produces exactly this dissociation, the decomposition that makes the explanations testable against each other, and a measured map of where the failure does and does not appear.

We audit the assumption end to end. The situation turns out worse, and more structural, than a bias–variance account suggests. Our claims come in three tiers:

1. **Possibility (proved).** One-sided corrections can flip the sign of a prompt’s gradient at arbitrarily small policy divergence (Section 3). Group normalization does not fix the sign. The metrics practitioners monitor—KL, ESS, gradient cosine—provably cannot detect it. One warning sign survives on stale data alone: the two one-sided cells disagreeing with each other.
2. **Regime map (measured).** Whether the flip actually happens depends on the task–capability regime (Section 5). We chart it with a pre-registered protocol and find three regimes: one where one-sided selection measurably loses precision, one where stale selection tracks the oracle, and one where selection degenerates outright.
3. **Impossibility (proved + measured).** Where selection matters most, neither the decision nor its evaluation can be secured at a practical budget. Boundary certification needs 5–8 orders of magnitude more rollouts than anyone spends (Section 6). And a measurement ceiling caps what any estimator can even *demonstrate* against a noisy oracle (Section 4).

Figure 1 summarizes the framework and how the sections fit together.

The central object of this paper is not an estimator but a measurement: the *split-half oracle floor*. The oracle itself is built from rollouts, so its own ranking carries noise. The floor quantifies that noise: split the oracle’s rollouts into two independent halves, rank the prompts with each half, and record the top- k agreement between the two rankings. A low floor means the pool offers little recoverable ranking signal at that rollout budget, and our ceiling result (Proposition 2) shows that no method can demonstrate more precision than the floor permits. We applied this standard to our own pools by running FIRST (Floor-Informed Reliability of SelecTion), the pre-registered diagnosis of Section 8, and obtained the verdict *structural absence*: on saturated or tie-degenerate pools, increasing the oracle budget up to $K=256$ rollouts per prompt is not predicted to create a signal regime. This verdict, rather than a new estimator, is the paper’s main empirical finding, and we argue it is the more useful one for practitioners deciding where to spend fresh rollouts.

One further observation applies to every experimental section. Selection evaluation tends to fail in the optimistic direction without producing any visible error. During this project our own pipeline generated four such evaluation artifacts, each of which initially looked like a publishable positive result, and later detected all four (Section 7). We document each mechanism together with its guard, because the same mechanisms can arise in any experiment that reuses rollouts for data selection.

2 Related work

Off-policy influence and data selection. Influence scores built on a single token’s importance ratio rank prompts from offline trajectories [CROPI]. In our decomposition this is g_{00} : neither axis is restored. The appendix of that work already reports the symptom that motivates our estimand: per-prompt gradients agree pointwise (cosine above 0.6 on 40 of 50 prompts), yet only 28.8% of the top-10% neighborhood structure is preserved. Corrections that reweight the visited prefix [A Step Back,




Figure 1: Framework overview. *Estimate*: stale-rollout training-data ranking uses one of four corrections in a 2×2 decomposition; one-sided cells can silently reverse rankings and the monitored metrics cannot see it (§3). *Measure*: the split-half floor of the oracle, computed under four guards, quantifies recoverable signal and induces a ceiling on what any method can demonstrate (§4). *Decide*: an artifact-only, pre-registered diagnosis routes each pool to one of three regimes; in ours the verdict was structural absence (§8).

[Cumulative](#)] restore the occupancy axis (g_{10}). Their guarantees, however, assume a target-policy advantage, and practical critic-free objectives replace that advantage with a behavior-generated outcome reward; this substitution is the gap our counterexamples use. Corrections that reweight the sampled continuation [[Multi-Step](#)] restore the outcome axis (g_{01}), but not occupancy, when used for per-prompt gradient ranking. Full-trajectory corrections [[TIC-GRPO](#)] restore both axes, but their variance grows with response length, so they serve as a cost reference rather than a practical selector. The same tension appeared for Data Shapley in supervised data selection: valuation scores failed to outperform random selection until the selector was redesigned [[Rethinking-Shapley](#), [NASH](#)]. Our protocol (§8.1) is the RLVR counterpart, with one difference: the repair comes from a reliability measurement, not from a new estimator.

Rollout allocation and learning stability. Allocation methods [[GradAlign](#), [LearnAlign](#)] decide where to spend fresh rollouts to maximize training yield, and stale-data stabilizers keep *learning* well-behaved under reuse. Neither certifies, or even measures, the correctness of the top- k *selection decision*, which is what practitioners actually consume. Classical off-policy correction and subset selection provide the statistical background: conditional importance sampling [[Rowland et al.](#)], stationary-distribution correction [[Liu et al.](#)], doubly-robust policy gradients [[Huang & Jiang](#)], and

information-complexity lower bounds for top- k identification [Kaufmann et al.]. Our certification result (Section 6) is an empirical instance of the last of these: when value gaps at the ranking boundary vanish, identification cost diverges. A general budgeted algorithm for top- k certification also exists: it concentrates expensive queries at the ranking boundary and assumes the cheap scores are valid up to noise [ACE]. Our results identify the case this assumption excludes. A stale one-sided score can be systematically biased (sign-flipped), and then certification fails through bias, not through budget. The two failure ingredients we study also appear in adjacent work on on-policy distillation [TIDE]. There, degenerate repetition loops drive token-level agreement to near-zero divergence while outputs collapse, so a pointwise metric misses the object being consumed; and raw token log-ratios are dominated by the extreme tail, with the top 1% of tokens carrying roughly half the gradient—the same heavy-tail behavior we measure as 0.1% ESS for the full correction.

Reliability of noisy rankings. Our floor and ceiling analysis imports split-half reliability and the Spearman–Brown prophecy [Spearman, Brown] into RLVR selection evaluation. Concurrent work applies the same split-half logic as a performance ceiling in a different domain (crowd reading) [Floor-Ceiling]; our contribution is therefore the transplant of this reliability logic into RLVR selection evaluation and its pre-registered self-application, not the ceiling idea itself. To our knowledge, no prior work in this literature reports the reliability of its own oracle, and Section 4 shows that several of our initially promising comparisons no longer hold once it is reported.

3 Theory: one-sided corrections and what the monitored metrics cannot see

3.1 A 2×2 decomposition of stale-gradient estimators

For prompt x , response $Y = (A_1, \dots, A_T)$ with terminal verifiable reward $R(Y)$, the current-policy gradient is $g_\pi(x) = \sum_t \mathbb{E}_{H_t, A_t \sim \pi} [Q_t^\pi(H_t, A_t) z_t]$ with $z_t = \nabla_\theta \log \pi(A_t | H_t)$. With behavior trajectories, token ratios $r_j = \pi(A_j | H_j) / \beta(A_j | H_j)$, prefix products $P_t = \prod_{j < t} r_j$ and suffix products $S_t = \prod_{j > t} r_j$, four population quantities arise from the same stale data:

$$\begin{aligned} g_{00} &= \sum_t \mathbb{E}_\beta [r_t R z_t] && (\text{occupancy } \beta, \text{ continuation } \beta) \\ g_{10} &= \sum_t \mathbb{E}_\beta [P_t r_t R z_t] && (\text{occupancy } \pi, \text{ continuation } \beta) \\ g_{01} &= \sum_t \mathbb{E}_\beta [r_t S_t R z_t] && (\text{occupancy } \beta, \text{ continuation } \pi) \\ g_{11} &= \sum_t \mathbb{E}_\beta [P_t r_t S_t R z_t] = g_\pi. \end{aligned}$$

The subscripts state which axis is restored. The exact error decompositions

$$g_{10} - g_\pi = \sum_t \mathbb{E}_{d_{\pi, \pi}} [(Q_t^\beta - Q_t^\pi) z_t] \quad (\text{continuation bias}), \quad g_{01} - g_\pi = \sum_t (\mathbb{E}_{d_{\beta, \pi}} - \mathbb{E}_{d_{\pi, \pi}}) [Q_t^\pi z_t] \quad (\text{occupancy bias})$$

show that the two one-sided estimators remove *different* biases; they are not substitutes.

3.2 Sign reversal at arbitrarily small divergence

Theorem 1 (One-sided sign reversal). *For every $\epsilon > 0$ there exist two-token environments and policy pairs (π, β) with trajectory KL $O(\epsilon^2)$ such that $g_\pi = -\epsilon/2$ while $g_{00} = g_{10} = +\epsilon/2$ (continuation counterexample), and likewise $g_{00} = g_{01} = +\epsilon/2$ (occupancy counterexample). Group normalization with any group size $K \geq 2$ preserves the reversal.*

The two constructions are given in Appendix A; both flip a Bernoulli success profile so that the policies agree wherever the retained correction looks and disagree exactly where the discarded half lives. Both are machine-checked (raw gradients, exact enumeration of group normalization for $K \in \{2, 4, 8\}$ and $\epsilon \in \{0.4, 0.1, 0.01\}$) by a self-contained verification script shipped with the code.

Corollary 1 (Metric blindness). *No bound on trajectory KL, no lower bound on ESS, and no norm-blind, scale-invariant single-estimator diagnostic (e.g. pointwise cosine of the estimator against itself across checkpoints) can exclude sign reversal of a prompt’s selection score: in the constructions above all such quantities are driven to their benign limits while the ranking direction is reversed. A gradient-norm floor or a selection-margin bound is necessary; and since $\epsilon \rightarrow 0$ also drives the absolute gradient to zero, normalized cosines conceal exactly the missing signal floor.*

Proposition 1 (Disagreement accompanies double failure). *In the continuation counterexample, $\cos(g_{10}, g_{01}) = -1$: the two one-sided cells disagree maximally.*

Empirically—an observation, not a theorem—in a seeded random search over 5×10^4 two-token, two-parameter environments we found *zero* instances of a double reversal (both one-sided estimators reversed against g_π) accompanied by cell agreement ($\cos(g_{10}, g_{01}) > 0.9$); the identity $g_{11} \equiv g_\pi$ was verified on every sample.

Proposition 1 is deliberately modest: absence of disagreement is *not* proven to imply safety in general architectures with parameter sharing, disagreement cannot certify *which* cell is right, and its practical resolution is limited by the same margin collapse quantified in Section 6. We therefore use cell disagreement only as an *unreliability signal*—a cheap, stale-computable reason to distrust a ranking—never as a selector. The mechanism is transparent from the error decompositions: g_{10} and g_{01} subtract different bias terms from the same truth, so a large simultaneous failure of both requires the two bias terms to agree, which the decomposition makes non-generic.

4 Measuring selection: floor, ceiling, and when comparison is meaningful

Estimand. All experiments target the decision practitioners consume: the top- k set by $s_i = \cos(\mu_i, v)$, where μ_i is the prompt’s mean current-policy gradient and v a validation mean gradient. We report top- k precision against an on-policy oracle, boundary margins, and rollout budgets. Pointwise cosine averages are diagnostics, never primary results.

The split-half floor. The oracle itself is a K -rollout Monte Carlo estimate. We therefore report, for every run, the *corrected split-half floor*: rollouts are split into two disjoint halves (odd/even micro-groups), each half induces a top- k set (ties broken by *independent* per-half jitter; Section 7 explains why this correction matters), and the floor is the mean overlap fraction over 20 jitter pairs. Chance is k/n . The floor measures recoverable ranking signal at budget K ; it is conservative (an estimator can measurably exceed it), and Section 7 shows it is itself vulnerable to tie artifacts unless computed exactly this way.

A measurement ceiling. The floor does more than describe the oracle: it bounds what any comparison against that oracle can show.

Proposition 2 (Measurement ceiling). *Model prompt scores as exchangeable Gaussians: true value t_i , half-oracle scores a_i, b_i with $\text{corr}(a, b) = \rho_h$, full- K oracle score o_i with reliability $\rho_f = 2\rho_h/(1+\rho_h)$*

Table 1: Floor $\rightarrow$ ceiling mapping under Proposition 2 ($n=512$, $k=51$, chance 0.100; simulation, 600–800 replicates/cell). The v2 weak-signal regime (floors 0.137–0.176) caps *any* method’s measurable precision at ≈ 0.30 – 0.37 ; the observed stale baseline is already at this ceiling.

observed floor	0.104	0.113	0.130	0.139	0.170	0.262	0.383
half-reliability ρ_h	0.02	0.05	0.10	0.12	0.20	0.40	0.60
ceiling on measurable precision	0.170	0.219	0.278	0.300	0.372	0.516	0.640

(Spearman–Brown). Then for any estimator s computed independently of the oracle’s sampling noise, the expected top- k overlap between s and the K -rollout oracle is at most that of the best possible estimator $s = t$, namely $\text{Ovl}(\sqrt{\rho_f})$, where $\text{Ovl}(c)$ is the expected top- k overlap of two Gaussian scores with correlation c . The observable floor equals $\text{Ovl}(\rho_h)$, so the ceiling is a computable, monotone function of the observed floor.

Table 1 gives the mapping at our experimental geometry ($n=512$, $k=51$), validated by simulation (Appendix B). Two consequences organize everything that follows. First, when the floor is near chance the ceiling collapses toward chance: *under the proposition’s model, no estimator can measurably outperform, and comparative evaluation of selection methods is void*. Second, an estimator observed at the ceiling has plausibly exhausted the measurable signal—further method development cannot show gains without first raising the floor (more oracle rollouts, or a pool with more signal).

Retention and its domain. When comparing regimes we report retention $(\text{prec} - c)/(\text{floor} - c)$ with c the chance level, declared *undefined* whenever $\text{floor} - c \leq 0.02$; saturated cells otherwise produce sign-flipping or exploding ratios (the 14B GSM8K cell below is the cautionary case).

5 The regime map: where the flip does and does not happen

Setup. Qwen2.5-7B/14B-Instruct; GSM8K, MATH-500, and DAPO-Math-17k; drift between β and π is controlled by LoRA RFT for 50–400 steps; gradients are JL-projected; oracle uses $K=32$ fresh rollouts per prompt (8 micro-groups $\times$ 4); pools are $n=256$ (v1) or $n=512$ (v2) with $k=10\%$. The v2 protocol was pre-registered after a full audit (equal per-cell budgets, seeded tie-breaks, independent jitter, manifests); judgment criteria were fixed before results were read. Two v2 seeds per task have completed at the time of writing; the decision rule is frozen, and the verdict will be recomputed under it when the remaining seeds arrive.

5.1 Where signal exists, one-sided selection can lose it

The v1 gate (7B GSM8K) passed its pre-registered criteria: one-sided estimators lose top-10% precision against the uncorrected baseline. At the largest drift (400 LoRA steps), the prefix-corrected estimator falls to precision 0.040 (overlap 1/25). The exact hypergeometric tail for this cell is $P(X \leq 1) = 0.27$, so we report it as a directional observation, not as an individually significant below-chance claim. One transparency note applies. The v1 floors were originally computed with the shared-jitter estimator that Section 7 later diagnosed as pitfall P2, which had inflated them to 0.32. The corrected re-estimation (Table 5) puts the drift-400 floor at 0.206. This is still a $2\times$ -chance signal regime, and the 0.040 collapse lies far below both floor and chance; in other words, even the regime most favorable to our thesis had to survive our own guards. Downstream, fine-tuned accuracy after selection (base 0.740) spans 0.820–0.880 across oracle-, one-sided-, and

Table 2: v2 corrected results ($n=512$, $k=51$, chance 0.100; 2 seeds). Floors are corrected split-half (independent jitter). DAPO precisions are jitter-sensitive (± 0.08) due to margin-0 tie degeneracy and shown as bands.

run	floor	g_{00}	g_{10}	g_{01}	g_{11}	regime
GSM8K s1	0.137	0.294	0.275	0.196	0.196	weak signal; stale at ceiling
GSM8K s2	0.176	0.235	0.196	0.216	0.216	weak signal
DAPO s1	0.176	0.04–0.16 (tie-degenerate)				margin 0.0; chance
DAPO s2	0.176	0.04–0.16 (tie-degenerate)				margin 0.0; chance

random-selected pools, with no consistent ordering. In this saturation-adjacent regime, misranking has little room to change downstream accuracy, which is consistent with the results below.

On 7B MATH-500 the ordering flips. This is the strongest and least noisy signal regime we measured: the corrected floor is 0.288, and its K' -curve *rises* (Table 5). Here the uncorrected g_{00} is worst (0.160), and every stale estimator retains only 0.28–0.64 of the available signal. At 14B MATH-500 the loss concentrates on the suffix axis (signal retention $g_{01}=0.44$ vs. $g_{10}=1.00$), although the declining K' -curve there makes the absolute floor less certain. Which cell fails is therefore regime-dependent. What does not change is this: *some cell fails while every monitored metric looks normal*—exactly what Theorem 1 permits, and no more.

5.2 Where signal is weak, stale selection tracks the oracle

The corrected v2 runs put every completed condition at floors 0.137–0.176 against chance 0.100 (Table 2). Here stale scores are *not* distinguishable from the oracle; they track it. On GSM8K, the uncorrected stale baseline reaches precision 0.294/0.235 (seeds 1/2), with upper-tail exact $p = 2.8 \times 10^{-5}$ and 2.1×10^{-3} , and Fisher-combined $p = 1.0 \times 10^{-6}$. (The two seeds share the prompt pool; conditional on the pool, the hypergeometric nulls are driven by independent rollout randomness, which is what the combination assumes.) The value 0.294 is also *at the measurement ceiling* (≈ 0.30 at floor 0.137; Table 1). The stale estimator has therefore reached the limit of what the oracle can resolve, so neither the losses of its one-sided competitors nor the gains of any other method are measurable in this regime. The DAPO pool is tie-degenerate: boundary margins are exactly 0.0, and precision itself moves by ± 0.08 under tie-break jitter, so we report bands. There, all stale estimators are at chance, and so is everything else.

[PENDING — v2 seed 3, running] Additional GSM8K/DAPO seeds are being harvested; Tables 2 and 5 will gain one row group each and all seed-pooled statistics will be recomputed under the pre-registered rules. Verdict logic is frozen; only numbers change. *Reserved space: half a page (updated tables + two sentences).*

Prompt-level reversal prevalence (preliminary). Top- k overlap is a coarse summary. A reviewer will ask a finer question: how often does the sign of an individual prompt actually flip? Table 3 answers this for the completed v2 GSM8K seeds. Stale scores disagree in sign with the fresh oracle on 36–48% of signal-bearing prompts overall, on 24–35% of prompts inside the decision band, and on 6–15 of the oracle’s top-51 prompts. These raw rates cannot be read against zero, because the oracle also disagrees with itself: an independent half of the oracle flips sign against the other half on 39%/43% of the same prompts (decision band: 27%/20%), as the reference row reports.

Table 3: Prompt-level sign reversal against the fresh oracle (completed v2 GSM8K seeds; zero-signal prompts excluded, denominators shown). Raw rates must be read against the reference row—the oracle’s split-half self-disagreement on the same prompts—not against zero. Decision band = oracle ranks $k \pm 26$.

	overall reversal		decision-band reversal	
	s1	s2	s1	s2
g_{00}	67/163	80/168	12/37	12/34
g_{10}	68/163	73/168	13/37	10/34
g_{01}	61/163	63/168	10/37	8/34
g_{11} (full)	59/163	64/168	9/37	8/34
oracle self-disagreement (reference)	66/168	77/180	10/37	7/35

Overall reversal for every cell lies within a few points of this reference. We therefore do not read the overall rates as estimator failure; they are dominated by oracle noise. The information lies in two contrasts that place the same oracle noise on both sides.

The first contrast concerns the decision band. There, the one-sided cells lie above the reference while the full correction lies on it: pooled band reversal is 0.34 (g_{00}) and 0.32 (g_{10}), against 0.24 for g_{11} and 0.24 for the reference, with g_{01} indistinguishable from full (0.25). Prompt-paired McNemar tests against g_{11} (exact, over all signal-bearing prompts) reach significance for g_{00} in seed 2 (27 vs. 11 discordant prompts, $p = 0.014$); the remaining paired contrasts are not significant ($p \geq 0.13$). Pooling is over a shared prompt pool, so we present these numbers as replication across seeds, not as independent evidence. This prompt-level resolution also exceeds what top- k precision offers in the same regime (Table 2), where the ceiling hides any separation.

The second contrast tests the theory’s warning sign at the prompt level. Conditioned on the two one-sided cells disagreeing in sign, the prefix-corrected cell g_{10} is the reversed one in 59%/62% of cases (seeds 1/2), against a 37% reversal rate under agreement (Fisher exact $p = 0.015/0.007$). The same pattern replicates in the strong-signal v1 condition (7B, MATH-500): 81% against 45%, $p = 0.012$. In the saturated 14B regime, where selection itself degenerates (§5.3), the warning is correspondingly uninformative ($p \geq 0.44$); this is expected, because no signal remains there to protect. These results support the disagreement signal of Proposition 1 in real LLM runs, with one estimand caveat stated explicitly. For scalar per-prompt scores, a *double* reversal appears as cell *agreement*; the 37–45% reversal rate under agreement is exactly the case this warning cannot detect. This is consistent with “agreement guarantees nothing,” and it is distinct from the proposition’s vector-cosine statement. The DAPO seeds contribute only 2–4 signal-bearing prompts each (tie degeneracy, §5) and are excluded.

5.3 Where capability saturates, selection degenerates

At 14B on GSM8K ($\sim 95\%$ accuracy), the pool contains 132/256 prompts with zero reward variance. Every estimator’s precision equals the chance level (0.087–0.111 vs. 0.098), and the oracle disagrees *with itself* below chance (floor $0.080 < 0.098$). Selection, ranking, and certification therefore degenerate together: no method and no certificate can have an effect here. The practical consequence is that selection pipelines must scale candidate difficulty with model capability; otherwise they spend, without noticing, on machinery that cannot change the outcome.

6 Certification is out of reach where it matters

Even where a point estimate selects well, one may ask for the decision to be *certified*: spend fresh rollouts near the ranking boundary until confidence intervals separate the top- k set. We implemented this (sequential boundary refinement with shared-validation error propagation; pseudocode in Appendix C)—not as a method contribution, but as a measuring instrument. The instrument returns a structural negative. Across every completed condition (both scales, three tasks, drift 50–400, selection fractions 5–25%), certification fails at budgets 1.02 – $1.04\times$ the uncorrected cost. The reason is geometric: top- k boundary margins are 0.00 – 0.24° , while achievable confidence radii are 51 – 180° . The gap is not a constant factor. With radius $\alpha_v(m) = \arcsin(c/\sqrt{m})$, reaching the observed margins requires 4.1×10^5 to 1.0×10^8 times more observations—5–8 orders of magnitude in budget. The gap is robust to (ϵ, δ) -PAC slack and to deepened validation. The cause is that prompt values are intrinsically dense at the selection boundary: *exact certification fails precisely where selection matters—at every margin we observed*. Approximate-membership and indifference-zone formulations are weaker demands our measurements do not exclude. Fresh rollouts spent on exact certification yielded no improvement in any condition we measured; Section 9 argues they belong in floor measurement instead.

7 Four evaluation pitfalls our own pipeline detected

Every mechanism below produced, at some point in this project, an apparently valid positive result that survived a first review and failed under a targeted audit. We report them as findings in their own right. Each comes with a guard now built into our protocol, and we suspect instances of each exist in the current rollout-reuse literature.

P1: Oracle-sample reuse (leakage). Our v1 hybrid intervention (completing the “missing half” with fresh continuations) showed a perfect 21/21-condition recovery, mean $+0.46/+0.49$ precision—which would have been the central causal result of the paper. A code audit found the treated cell reused rollouts that also defined the oracle; the shared noise, not causation, produced the recovery. A pre-registered equal-budget decisive experiment with disjoint samples showed *no consistent recovery* (e.g. GSM8K s1: fully-fresh cell 0.38 *below* mixed cell 0.46). *Guard*: strict sample disjointness between treatment and oracle (odd/even micro-group separation), enforced structurally.

P2: Shared-jitter tie inflation. Our first floor implementation broke score ties with a jitter shared across the two halves. On a tie-degenerate pool (DAPO, margins 0.0) this inflated the floor from 0.176 to 0.804, which briefly supported an incorrect “inverse correlation” conclusion. The same estimator had also inflated the v1 floors (0.32 reported vs. 0.15–0.21 corrected), which means every absolute floor in this literature computed with tie-unaware agreement warrants re-examination. *Guard*: independent per-half jitter streams, 20 seeded pairs, plus reporting precision *bands* under jitter on tie-heavy pools.

P3: Cell-budget imbalance. Comparing estimator cells with unequal effective rollout counts favors the cell with the larger budget without any visible sign of it; our v1 hybrid grid had this flaw in addition to P1. *Guard*: equal- K redesign of all cell comparisons, verified from manifests.

P4: Saturation-shared overlap inflation. An above-chance floor can be an artifact of *shared liveness* rather than shared ranking: prompts that saturate in both halves (all-correct or all-wrong)

Table 4: K -scaling floor curves (observed exact recomputation; prediction by Spearman–Brown + Gaussian overlap). Chance 0.100. Note the *decreasing* observed curves—the P4 artifact signature—and the near-zero half-reliabilities. Verdict (pre-registered): structural absence.

run	r_1	$K'=8$	16	24	32	pred. 128	pred. 256
GSM8K s1	0.045	0.263	0.208	0.201	0.182	0.288	0.391
GSM8K s2	−0.005	0.220	0.182	0.170	0.156	0.094	0.094
DAPO s1	0.009	0.112	0.123	0.139	0.146	0.147	0.182
DAPO s2	−0.038	0.119	0.148	0.158	0.158	0.094	0.094

drop to an exactly-tied score, and the surviving “live” sets overlap between halves, producing top- k agreement with zero true ranking signal. In a synthetic pool with *no* ranking signal, saturation sharing alone yields floor 0.126 at $K'=8$, decaying toward chance (0.110) at $K'=32$. The signature of this artifact is a floor that *decreases* as the oracle budget grows. Our measured K -scaling curves (Section 8) show exactly this decreasing shape, which means even our reported weak-signal floors are best read as upper bounds. *Guard*: always report the floor as a function of K' ; treat a non-increasing floor curve as artifact evidence, and report tie mass alongside.

8 FIRST: diagnosing signal absence from existing artifacts

The pitfalls motivate the final contribution: FIRST, a pre-registered, artifact-only procedure (no new rollouts) that decides whether a pool supports selection *before* any new experiment is funded. It has three stages, all computed from artifacts every RLVR selection run already produces.

Stage 1: conditional-subset precheck. Any proposed pool intervention definable as a filter on existing prompts (e.g. “keep only live prompts, $0 < \text{pass-rate} < 1$ ”) already exists as a subset of completed runs. Computing the corrected conditional floor on that subset (with bootstrap CIs) prechecks the intervention at zero GPU cost: if the conditional floor does not clear $2\times$ chance in a majority of eligible runs, building the new pool has no evidential basis. On our pools this returned NO-GO, and the intervention (a live-only GSM8K pool) was cancelled without spending a GPU-hour.

Stage 2: K -scaling floor curve. From stored micro-group gradients the floor is *exactly* recomputed at $K'=8, 16, 24, 32$, and extrapolated to $K'=64\text{--}256$ by Spearman–Brown on the half-reliability r_1 mapped through the Gaussian overlap function, with a calibration check at the largest observed K' . The pre-registered verdict rule: if a majority of eligible runs are predicted to reach $2\times$ chance by $K' \leq 256$, recommend oracle expansion; otherwise declare structural absence.

Table 4 shows the result: one GSM8K seed extrapolates to $2\times$ chance at $K'=128$ from $r_1=0.045$; the other is flat at chance; both DAPO runs never reach it. Majority not met: *structural absence*—meaning no recoverable ranking signal *within the evaluated pools and budget range*. This is an epistemic verdict about what these pools and budgets can support, not an ontological claim that prompt-value differences do not exist.

Running the same analysis over every artifact-complete run in the project (twelve conditions including the v1 gate family; Table 5) widens the evidence base without changing the pre-registered vote. Signal regimes are rare (2–3 of 12). The two credible ones are exactly the conditions where the regime map’s one-sided losses occur (7B MATH-500 and maximum-drift GSM8K). And half of

Table 5: Extended K -scaling evidence over all twelve completed conditions (corrected floors; “shape” is the observed $K'=8 \rightarrow 32$ direction, with $\downarrow$ marking the P4 artifact signature). The pre-registered verdict is computed only on the v2 rows and is unchanged; the extension shows how rare signal regimes are: three conditions clear $2\times$ chance at $K'=32$, and only two of those with a rising curve. Floors here use random half-assignment resampling and may differ slightly from the fixed odd/even-split floors of Table 2.

condition	task	r_1	floor @ $K'=32$	reaches $2\times$ chance	shape
v1 drift-50	GSM8K	0.020	0.177	pred. $K'=256$	$\uparrow$
v1 drift-100	GSM8K	0.053	0.146	pred. $K'=128$	$\downarrow$
v1 drift-200	GSM8K	0.067	0.147	pred. $K'=64$	$\downarrow$
v1 drift-400	GSM8K	0.065	0.206	at $K'=32$	$\uparrow$
v1 14B	GSM8K	0.018	0.113	never	$\downarrow$
v1 7B	MATH-500	0.070	0.288	at $K'=32$	$\uparrow$
v1 14B	MATH-500	0.075	0.221	at $K'=32$	$\downarrow$
v2 s1/s2	GSM8K	0.045/−0.005	0.182/0.156	pred. 128/never	$\downarrow\downarrow$
v2 s1/s2	DAPO	0.009/−0.038	0.146/0.158	never	$\uparrow\uparrow$
v2 s2	MATH-500	0.001	0.181	never	$\downarrow$

all observed curves *decline* in K' —the P4 artifact signature—so absolute floors elsewhere are best read as upper bounds.

Stage 3 (replay): the ceiling made visible in budget space. As a final check we replayed all fresh-budget allocation policies on the completed v2 runs against a sample-disjoint truth (odd micro-groups only): stale cells, a pass-rate heuristic, random, uniform and boundary audits at 1–25% budgets, a sequential disagreement-guided audit, full-fresh at up to 100% budget, and our own pre-registered floor-gated protocol. The result is the measurement ceiling made concrete. On weak-signal GSM8K, no fresh budget yields an improvement. Full-fresh at 100% budget scores *lower* (0.167) than the zero-cost stale baseline (0.235). Our own pre-registered floor-gated protocol also loses (0.172); by its pre-registered loss condition, it is withdrawn from prescriptive use and kept only in the diagnostic role this section describes. The best zero-cost policy is not a gradient method at all but the behavior pass-rate heuristic (0.255). On DAPO, every estimator is at chance (0.08–0.16, truth margins 0.0). Throughout these replays, the monitored metrics stay in their normal ranges, precisely as Corollary 1 warns. The full correction’s effective sample size is 0.1% (clip fraction 0.94–1.00); this is the variance-side reason one-sided truncations exist in the first place, measured in the same runs that show what the truncation costs.

Two pre-registered caveats accompany the method. First, the calibration check shows that the Gaussian model under-predicts observed floors at $K'=32$ (e.g. 0.094 predicted vs. 0.156 observed). This is consistent with the P4 artifact inflating the *observed* values, and it makes the verdict conservative in the safe direction. Second, the single positive extrapolation rests on an r_1 small enough to be affected by the P4 artifact, so we treat it as motivation for an inexpensive re-check when the remaining seeds arrive, not as grounds for a GPU expansion.

8.1 From diagnosis to decision: the FIRST protocol

The verdict is not the end of the pipeline. It is the first branch of a decision. Each regime in the map comes with a selection policy, and our own replays already measure the cost of each one. The

diagnosis therefore extends into a decision protocol. The data-valuation literature encountered the same problem in supervised learning: when Shapley-based selection proved no better than random, the response was a corrected selection method, not abandonment of selection [Rethinking-Shapley, NASH].

The FIRST protocol inputs: the behavior rollouts the run already has (optionally a 48-prompt fresh probe); fix the threshold $\tau = 2 \times \text{chance}$ *before* looking.

1. **Measure.** Compute the corrected split-half floor (independent jitter) and the K -scaling curve with its Spearman–Brown budget estimate K^* (§8); run the four pitfall guards (§7).
2. **Branch.**
 - (a) *No reachable signal* (floor $\approx$ chance, no finite K^*): abstain—sample uniformly and spend nothing on scoring. Every estimator we measured is at chance here, and the oracle disagrees with itself (§5.3).
 - (b) *Weak signal* (chance $<$ floor $<$ τ): rank by the zero-cost behavior pass rate. In our replays it matched or outperformed every gradient estimator (0.255 vs. 0.235 best stale), and the measurement ceiling shows no method can *measurably* exceed it in this regime (§4).
 - (c) *Measurable signal* (floor $\geq \tau$): stale scores are usable. Compute both one-sided cells anyway, and send the prompts where they disagree in sign to a fresh audit (or drop them). Under disagreement the reversal rate was 59–81%; under agreement, 37–45% (three runs, $p \leq 0.015$). Agreement guarantees nothing—this is the case the warning cannot detect, and it is stated as such.
3. **Report.** Attach the floor, the ceiling, and the branch taken to any downstream selection claim; a result above its own oracle’s ceiling is an artifact by Proposition 2.

The verdict costs no new rollouts; branch (c)’s audit spends fresh budget only on the disagreement set, which in our runs was 14–25% of signal-bearing prompts.

[PENDING — **capability–difficulty sweep (decision gate 9/10)**] A small-model sweep (Qwen3.5 0.8B/2B/4B on the same pools) would chart the floor as a function of the capability–difficulty gap—the one controllable knob that could *create* a signal regime and turn the regime map into a causal phase diagram. If run, it adds one figure (floor phase map) and one paragraph here; if not, the map remains observational, as stated in Limitations. *Reserved space: one figure + half a page.*

9 Practical guidance, limitations, and outlook

For practitioners. (1) *Measure the floor first*—run FIRST and take the branch its protocol prescribes (§8.1). A 48-prompt fresh probe is enough to estimate the split-half floor. Below $2 \times$ chance (fix the threshold *before* looking), do not spend on selection. In our replays the best weak-signal policy was the zero-cost pass-rate heuristic, random was close behind, and no fresh budget up to 100% produced an improvement. (2) *Never budget for boundary certification*; the 5–8 order-of-magnitude gap is structural. (3) *Scale candidate difficulty with capability*, or selection degenerates regardless of estimator. (4) *Report floor and ceiling with any selection comparison*;

a claimed improvement above the measurement ceiling of its own oracle indicates assumption violation, estimator–oracle dependence, or an evaluation artifact. (5) *Distrust one-sided rankings when the two one-sided cells disagree*; agreement guarantees nothing, but disagreement is a zero-cost warning—formalized as branch (c) of the protocol.

Limitations. Our empirical evidence covers Qwen2.5 (7B/14B), math tasks, and LoRA drift, with two completed v2 seeds; the remaining seeds and a small-model capability–difficulty sweep are running, and the pre-registered verdict logic will be re-applied when they arrive. The ceiling proposition assumes an exchangeable Gaussian score model; simulation validates it at our geometry, and a distribution-free proof remains open. The disagreement result is an existence-plus-search finding, not a general sufficiency theorem. The theory counterexamples are two-token constructions by design: they show what the monitored metrics *cannot* exclude, not what typically happens, and the regime map is their empirical complement.

Safety relevance. Data selection is a curation channel: whatever a misranked estimator systematically promotes is amplified by subsequent training. Corollary 1 therefore has a safety reading. If stale selection promotes prompts on which shortcut or reward-hacking behavior scores well, the drift metrics a team monitors cannot flag the promotion, for the same reason they cannot flag the sign reversal itself. The selection layer is then an unaudited input to the reward loop. The measurement-first protocol applies unchanged when the selection target is a safety metric rather than accuracy, and we recommend running it there before trusting any stale-scored curation.

Outlook. The negative results are central rather than incidental: they delimit where the field’s growing investment in stale-rollout data selection can succeed (unsaturated, tie-free, signal-bearing pools—which our diagnosis identifies at no extra cost) and where it cannot. The measurement-first protocol—floor, ceiling, preregistration, and the four guards—is applicable to any selection-method paper, including ours.

References

- [CROPI] Data-Efficient RLVR via Off-Policy Influence Guidance. ACL 2026. <https://arxiv.org/abs/2510.26491>
- [A Step Back] A Step Back: Prefix Importance Ratio Stabilizes Policy Optimization. <https://arxiv.org/abs/2601.22718>
- [Cumulative] Rethinking Importance Sampling in LLM Policy Optimization: A Cumulative Token Perspective. <https://arxiv.org/abs/2605.07331>
- [Multi-Step] Multi-Step Likelihood-Ratio Correction for RLVR. <https://arxiv.org/abs/2605.20865>
- [TIC-GRPO] TIC-GRPO: Trajectory Importance Correction for GRPO. <https://arxiv.org/abs/2508.02833>
- [GradAlign] GradAlign. COLM 2026. <https://arxiv.org/abs/2602.21492>
- [TIDE] Mismatch Matters: On-Policy Distillation Beyond Token Agreement. <https://arxiv.org/abs/2608.09836>

- [LearnAlign] LearnAlign. <https://arxiv.org/abs/2506.11480>
- [Rowland et al.] Rowland et al. Conditional Importance Sampling for Off-Policy Learning. AISTATS 2020.
- [Liu et al.] Liu et al. Off-Policy Policy Gradient with Stationary Distribution Correction. UAI 2019.
- [Huang & Jiang] Huang & Jiang. From Importance Sampling to Doubly Robust Policy Gradient. ICML 2020.
- [Kaufmann et al.] Kaufmann & Kalyanakrishnan. Information Complexity in Bandit Subset Selection. COLT 2013.
- [Spearman] Spearman. Correlation calculated from faulty data. 1910.
- [Brown] Brown. Some experimental results in the correlation of mental abilities. 1910.
- [ACE] Top- k on a Budget: Adaptive Certification at the Boundary. <https://arxiv.org/abs/2601.20989>
- [Floor-Ceiling] Floor, Ceiling, Fusion Gap: Split-Half Oracles for Crowd Reading. <https://arxiv.org/abs/2608.01704>
- [Rethinking-Shapley] Rethinking Data Shapley for Data Selection Tasks: Misleads and Merits. ICML 2024. <https://arxiv.org/abs/2405.03875>
- [NASH] Is Data Shapley Not Better than Random in Data Selection? Ask NASH. ICML 2026. <https://arxiv.org/abs/2605.10684>

A Counterexample constructions and machine verification

Continuation counterexample (reverses g_{00}, g_{10}). Two-token MDP; first action $A_1 \sim \text{Bernoulli}(q)$, $q=1/2$, identical under both policies; reward $R = A_2$. Behavior success profile $Q^\beta(A_1=0) = \frac{1}{2} - \epsilon$, $Q^\beta(A_1=1) = \frac{1}{2} + \epsilon$; current policy swaps the two values. The gradient w.r.t. the logit of q is $g_\pi = -\epsilon/2$ while $g_{00} = g_{10} = +\epsilon/2$: the prefix already matches perfectly, so a perfect prefix ratio changes nothing, and the behavior-policy *ending* drives the sign. Mean suffix KL is $O(\epsilon^2)$; the 1-D cosine is -1 throughout.

Occupancy counterexample (reverses g_{00}, g_{01}). Behavior visits $P(A_1=0) = \frac{1}{2} + \epsilon$, current $\frac{1}{2} - \epsilon$; the second action is $\text{Bernoulli}(1/2)$ under both, and $R = A_2$ if $A_1=0$, else $1 - A_2$. For the second-action logit, $g_\pi = -\epsilon/2$ while $g_{00} = g_{01} = +\epsilon/2$: no tokens remain after A_2 , so the future ratio is exactly 1, and the stale *visit* distribution drives the sign. Trajectory KL is again $O(\epsilon^2)$.

Group normalization. Both constructions have overall success rate $1/2$ under both policies; binary group standardization multiplies each direction by the same nonnegative scalar, hence preserves the reversal (exact enumeration for $K \in \{2, 4, 8\}$, $\epsilon \in \{0.4, 0.1, 0.01\}$ in the verification script, alongside leave-one-out unbiasedness and the confidence radius bound used in Section 6).

B Measurement ceiling: derivation sketch and validation

Under the model of Proposition 2, write the full- K oracle score as $o = \sqrt{\rho_f} t + \sqrt{1 - \rho_f} \eta$ with η independent noise. For any estimator s independent of η , $\text{corr}(s, o) = \text{corr}(s, t) \sqrt{\rho_f} \leq \sqrt{\rho_f}$, with equality iff s is (a monotone transform of) t . Expected top- k overlap between two jointly Gaussian scores is increasing in their correlation, which yields the ceiling $\text{Ovl}(\sqrt{\rho_f})$; the observable floor is $\text{Ovl}(\rho_h)$ by the same function. Simulation at $n=512$, $k=51$ (600–800 replicates per cell; independent-jitter tie handling) produces Table 1; the mapping is monotone and smooth, so inverting the observed floor through it is well-posed. The Gaussian assumption matters: heavy-tailed score distributions with a stable head can exceed the Gaussian ceiling estimate, which is precisely the P4 caveat and why we report the ceiling alongside, not instead of, the K -scaling curve.

C Sequential boundary certification (measuring instrument)

Off-policy scores serve as a prior; every candidate receives a small fresh micro-group; scores carry confidence balls (shared validation error propagated); fresh rollouts are then spent only on candidates whose balls straddle the top- k boundary, until $\min L_{\text{selected}} > \max U_{\text{rest}}$ or the budget is exhausted, in which case *certification failure is reported, never hidden*. Variants (scalar CI, (ϵ, δ) -PAC with $\epsilon \leq 0.05$) behave identically in all our conditions because the binding constraint is the boundary margin itself (Section 6).

D Preregistration and audit trail

All v2 judgment criteria, the precheck GO/NO-GO rule, the K -scaling verdict rule, and the reframing decision tree (“structural absence $\Rightarrow$ this paper’s structure”) were registered in the project log before the corresponding results were computed; the four pitfalls each carry a before/after commit pair. The verification script, manifests, and all readouts ship with the code release.